

Mars-JEPA: Multispectral Joint Embedding Predictive Architectures for Martian Landslide Segmentation

BMVC 2026 Submission # 1133

Abstract

The expansion of autonomous exploration on Mars has unlocked unprecedented scientific frontiers, yet it requires navigating a landscape characterized by frequent landslides that drive geomorphological change by mass movement of rock and debris. This presents critical mechanical hazards to robotic assets. Mitigating these risks through accurate landslide prediction is therefore essential to safeguarding mission longevity and preventing catastrophic loss. While landslide prediction from multispectral images using traditional methods yields high accuracy, its generalizability is limited. In this work, we study the performance of self-supervised learning-based (SSL) image models, specifically trained as joint-embedding predictive architectures (JEPA) and develop a novel physics-aware image encoder, showing additional performance gains while maintaining the generalizability provided by SSL models. Experiments reveal that our physics-aware masking strategy captures long-range dependencies between spectral bands. We show that by learning cross-spectral latent correlations rather than raw pixel reconstructions, our model is significantly more robust to the sensor noise and varying lighting conditions commonly encountered in multispectral remote sensing data. We train JEPA encoders employing different masking strategies, and benchmark their downstream landslide segmentation performance on the MMLSv2 dataset. This work directly improves autonomous safety for planetary exploration and establishes a scalable framework for self-supervised learning in high-dimensional remote sensing and terrestrial disaster response.

1 Introduction

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetur adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper, leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellentesque, augue quis sagittis posuere, turpis lacus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla facilisi. Sed a turpis eu lacus commodo facilisis. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetur.

Suspendisse vel felis. Ut lorem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, molestie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque vehicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.

Sed commodo posuere pede. Mauris ut est. Ut quis purus. Sed ac odio. Sed vehicula hendrerit sem. Duis non odio. Morbi ut dui. Sed accumsan risus eget odio. In hac habitasse platea dictumst. Pellentesque non elit. Fusce sed justo eu urna porta tincidunt. Mauris felis odio, sollicitudin sed, volutpat a, ornare ac, erat. Morbi quis dolor. Donec pellentesque, erat ac sagittis semper, nunc dui lobortis purus, quis congue purus metus ultricies tellus. Proin et quam. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Praesent sapien turpis, fermentum vel, eleifend faucibus, vehicula eu, lacus.

2 Related Work

The development of a robust landslide detection system for the Martian surface requires an interdisciplinary understanding of computer vision and planetary science. This section reviews recent Section 2.2, the specific challenges of Section 2.1, and the paradigm shift toward Sections 2.3 and 2.4.

2.1 Planetary Surface Analysis and Exploration

Mars has long been a focal point for planetary exploration, with significant research dedicated to identifying landslides on its surface. Detecting landslides automatically is a non-trivial task due to shifting atmospheric conditions and complex geomorphology. Historically, experts have identified these features manually using high-resolution imagery from the High Resolution Imaging Science Experiment (HiRISE) [1].

Early automated systems primarily utilized physics-based features derived from digital elevation models, such as slope and elevation. Recently, the field has transitioned towards deep learning, employing Convolutional Neural Networks (CNNs) [2] for automated mapping. While traditional studies relied on monochromatic imagery, the introduction of the MMLSv2 dataset [3] enabled the use of multimodal data, combining spectral imagery with topographic information to improve detection accuracy.

2.2 Advancements in Semantic Segmentation

Semantic Segmentation has advanced from basic pixel-wise classifiers to encoder-decoder frameworks designed to capture multi-scale context. Early designs like U-Net [4] became highly popular because of its simple design based on skip connections that preserved spatial features. While U-Net++ [5] improved this by using nested and dense skip pathways to bridge the semantic gap between the features. To handle objects at different scales, Feature Pyramid Network (FPN) [6] and Pyramidal Attention Networks (PAN) [7] focussed on integrating high-level global semantics with low-level spatial features.

Further refinements led to the development of modules like DeepLabv3+ [8] which relied on atrous convolutions to expand receptive fields. More recently attention driven architectures like MANet [9] and UPerNet [10] have sought to capture global dependencies effectively. Most recently transformer based frameworks like SegFormer [11] showed the efficiency of hierarchial self-attention in capturing long range concept.

2.3 Self-Supervised Learning in Remote Sensing

Self-supervised representation learning [12] has seen a paradigm shift from contrastive frameworks, which rely heavily on data augmentations, to Masked Image Modeling (MIM) [13]. While generative approaches like Masked Autoencoders (MAE) [14] have shown promising results by reconstructing raw pixels, they are often ill-suited for the multi-sensor nature of Martian remote sensing. Pixel reconstruction is computationally intensive for high-dimensional data and risks "information leakage" in multimodal datasets, where highly correlated channels (e.g., RGB and Grayscale) allow the model to bypass learning deeper geological features.

2.4 Joint Embedding Predictive Architectures (JEPA)

The JEPA framework has recently been proposed as a foundational step toward non-generative alternative for developing world models [1]. Unlike generative models, JEPA performs its predictive objective entirely within an abstract representation space, capturing. Following the I-JEPA approach [2], our model utilizes an asymmetrical structure consisting of a context encoder, a target encoder, and a predictor.

The context encoder processes visible portions of an image, while the predictor is tasked with estimating the latent embeddings of missing regions. This "prediction in latent space" encourages the network to capture underlying structural and semantic properties. A significant challenge in such self-supervised frameworks is representation collapse—where the encoder converges to a trivial constant output. To maintain feature diversity, methods such as Stochastic Information Gradient Regularization (SIGReg) [3], VICReg [4] are often employed to ensure that learned representations remain informative across multiple input dimensions.

3 Methodology

In this work, we study the downstream performance on segmentation tasks for two JEPAs trained with different tokenization masking strategies, on multispectral images. This section details the formulation and development of the self-supervised models for segmentation as follows: Section 3.1 formulates the landslide segmentation problem on multispectral images; Section 3.2 provides an overview of the general architecture and methodology to train a JEPA encoder; Section 3.3 discusses the adaptation of the I-JEPA [2] architecture for multispectral images; Section 3.4 provides a novel spatial-spectral masking strategy for training JEPAs on multispectral images, and finally Section 3.5 discusses the integration of the JEPA encoders with SOTA segmentation architectures for landslide segmentation on the Martian surface.

3.1 Problem Formulation

Consider a multispectral image $I \in \mathbb{R}^{C \times H \times W}$, where C is the number of bands, H, W are the height and width of the image respectively. We define a binary segmentation mask $M \in \{0, 1\}^{H \times W}$ over the pixels in the image, where a pixel $M[i, j]$ at row i and column j in the segmentation mask M has a value of 1 if it is predicted to be affected by a landslide, and 0 otherwise.

The Martian landslide detection problem is to develop a segmentation model $S_\theta : \mathbb{R}^{C \times H \times W} \rightarrow \{0, 1\}^{H \times W}$, parameterized by θ , to satisfy the objective

$$\max_{\theta} \mathbb{E}_{(I, M) \sim D} [\sigma_{\text{sim}}(S_\theta(I), M)], \quad (1)$$

where the multispectral image and corresponding segmentation mask pairs (I, M) are sampled from the dataset D , and $\sigma_{\text{sim}}(M_1, M_2) \in \mathbb{R}$ is a similarity metric between two binary segmentation masks M_1 and M_2 . Concretely, the objective aims to maximize the expected similarity between the segmentation mask predicted by the model and the actual segmentation mask, for all examples in the dataset, by performing an optimization over the model parameters θ .

3.2 Joint Embedding Predictive Architectures

3.3 Adapting I-JEPA for Multispectral Image Encoding

3.4 Spatio-Spectral JEPA

3.5 Segmentation Models with Self-Supervised Backbones

4 Experimental Setup

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetur adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper, leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellentesque, augue quis sagittis posuere, turpis lacus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla facilisi. Sed a turpis eu lacus commodo facilisis. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas

wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetur.

Suspendisse vel felis. Ut lorem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, molestie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque vehicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.

Sed commodo posuere pede. Mauris ut est. Ut quis purus. Sed ac odio. Sed vehicula hendrerit sem. Duis non odio. Morbi ut dui. Sed accumsan risus eget odio. In hac habitasse platea dictumst. Pellentesque non elit. Fusce sed justo eu urna porta tincidunt. Mauris felis odio, sollicitudin sed, volutpat a, ornare ac, erat. Morbi quis dolor. Donec pellentesque, erat ac sagittis semper, nunc dui lobortis purus, quis congue purus metus ultricies tellus. Proin et quam. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Praesent sapien turpis, fermentum vel, eleifend faucibus, vehicula eu, lacus.

5 Results and Discussion

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultrices. Lorem ipsum dolor sit amet, consectetur adipiscing elit. In hac habitasse platea dictumst. Integer tempus convallis augue. Etiam facilisis. Nunc elementum fermentum wisi. Aenean placerat. Ut imperdiet, enim sed gravida sollicitudin, felis odio placerat quam, ac

Encoder	Segmentation	MIoU	Precision	Recall	F1 Score
I-JEPA	UPerNet	0.8350	0.8613	0.8678	0.8643
	UNet	0.849	0.873	0.892	0.882
	UNet++	0.8247	0.8601	0.8795	0.8690
	FPN	0.8246	0.8488	0.8897	0.8687
	MANet	0.8100	0.8448	0.8683	0.8555
	DeepLabv3+	0.8183	0.8576	0.8689	0.8632
	PAN	0.8141	0.8559	0.8693	0.8622
SS-JEPA	UPerNet	0.865	0.900	0.905	0.902
	UNet	0.8521	0.8631	0.8868	0.8748
	UNet++	0.8399	0.8768	0.8596	0.8681
	FPN	0.8265	0.8656	0.8777	0.8716
	MANet	0.8262	0.8691	0.8703	0.8697
	DeepLabv3+	0.8202	0.8579	0.8660	0.8619
	PAN	0.8170	0.8620	0.8587	0.8604

Table 1: Segmentation performance with self-supervised encoder backbones.

Add citations for segmentation heads

Add IOU (FG) and IOU (BG) columns here.

pulvinar elit purus eget enim. Nunc vitae tortor. Proin tempus nibh sit amet nisl. Vivamus quis tortor vitae risus porta vehicula.

Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper, leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellentesque, augue quis sagittis posuere, turpis lacus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla facilisi. Sed a turpis eu lacus commodo facilisis. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetuer.

Suspendisse vel felis. Ut lorem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, molestie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque vehicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.

Sed commodo posuere pede. Mauris ut est. Ut quis purus. Sed ac odio. Sed vehicula hendrerit sem. Duis non odio. Morbi ut dui. Sed accumsan risus eget odio. In hac habitasse platea dictumst. Pellentesque non elit. Fusce sed justo eu urna porta tincidunt. Mauris felis odio, sollicitudin sed, volutpat a, ornare ac, erat. Morbi quis dolor. Donec pellentesque, erat ac sagittis semper, nunc dui lobortis purus, quis congue purus metus ultricies tellus. Proin et quam. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Praesent sapien turpis, fermentum vel, eleifend faucibus, vehicula eu, lacus.4

Model	MIoU	Precision	Recall	F1 Score
U-Net	0.814	0.858	0.868	0.863
U-Net++	0.823	0.864	0.879	0.871
PSPNet	0.828	0.866	0.884	0.875
DeepLabV3	0.817	0.870	0.860	0.865
DeepLabV3+	0.829	0.863	0.889	0.876
SegFormer	0.812	0.859	0.863	0.861
DualSwinV2 + U-Net++, 5-fold ensemble	0.867	0.900	0.911	0.905
DualSwinV2-S + Hybrid Decoder	0.828	0.856	0.899	0.877
SS-JEPA + UPerNet	0.865	0.900	0.905	0.902
I-JEPA + U-Net	0.849	0.873	0.892	0.882

Table 2: Generated by Spread-LaTeX

Add IOU (FG) and IOU (BG) columns here.

6 Conclusion

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Correct citations. Do not use Gemini for citations.

References

[1] Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with a joint-embedding predictive architecture, 2023. URL <https://arxiv.org/abs/2301.08243>.

[2] Randall Balestriero and Yann LeCun. Lejepa: Provable and scalable self-supervised learning without the heuristics, 2025. URL <https://arxiv.org/abs/2511.08544>.

[3] Adrien Bardes, Jean Ponce, and Yann LeCun. Vicreg: Variance-invariance-covariance regularization for self-supervised learning, 2022. URL <https://arxiv.org/abs/2105.04906>.

[4] Liang-Chieh Chen, Yukun Zhu, George Papandreou, Florian Schroff, and Hartwig Adam. Encoder-decoder with atrous separable convolution for semantic image segmentation. *CoRR*, abs/1802.02611, 2018. URL <http://arxiv.org/abs/1802.02611>.

[5] Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey E. Hinton. A simple framework for contrastive learning of visual representations. *CoRR*, abs/2002.05709, 2020. URL <https://arxiv.org/abs/2002.05709>.

- [6] Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked autoencoders are scalable vision learners. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 16000–16009, 2022.
- [7] Vlad Hondru, Florinel Alin Croitoru, Shervin Minaee, Radu Tudor Ionescu, and Nicu Sebe. Masked image modeling: A survey, 2025. URL <https://arxiv.org/abs/2408.06687>.
- [8] Yann LeCun and Courant. A path towards autonomous machine intelligence. 2022. URL <https://api.semanticscholar.org/CorpusID:251881108>.
- [9] Hanchao Li, Pengfei Xiong, Jie An, and Lingxue Wang. Pyramid attention network for semantic segmentation. *CoRR*, abs/1805.10180, 2018. URL <http://arxiv.org/abs/1805.10180>.
- [10] Rui Li, Shunyi Zheng, Ce Zhang, Chenxi Duan, Jianlin Su, Libo Wang, and Peter M. Atkinson. Multiattention network for semantic segmentation of fine-resolution remote sensing images. *IEEE Transactions on Geoscience and Remote Sensing*, 60: 1–13, 2022. ISSN 1558-0644. doi: 10.1109/tgrs.2021.3093977. URL <http://dx.doi.org/10.1109/TGRS.2021.3093977>.
- [11] Tsung-Yi Lin, Piotr Dollár, Ross B. Girshick, Kaiming He, Bharath Hariharan, and Serge J. Belongie. Feature pyramid networks for object detection. *CoRR*, abs/1612.03144, 2016. URL <http://arxiv.org/abs/1612.03144>.
- [12] Alfred S McEwen et al. Mars reconnaissance orbiter’s high resolution imaging science experiment (hirise). *Journal of Geophysical Research: Planets*, 112(E5), 2007.
- [13] Keiron O’Shea and Ryan Nash. An introduction to convolutional neural networks, 2015. URL <https://arxiv.org/abs/1511.08458>.
- [14] Sidike Paheding, Abel Reyes-Angulo, Leo Thomas Ramos, Angel D. Sappa, A. Rajaneesh, P. B. Hiral, K. S. Sajin Kumar, and Thomas Oommen. Mmlsv2: A multimodal dataset for martian landslide detection in remote sensing imagery. In *Proceedings of IEEE/CVF Computer Vision and Pattern Recognition (CVPR) Workshops*, 2026.
- [15] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *MICCAI*, 2015.
- [16] Tete Xiao, Yingcheng Liu, Bolei Zhou, Yuning Jiang, and Jian Sun. Unified perceptual parsing for scene understanding. *CoRR*, abs/1807.10221, 2018. URL <http://arxiv.org/abs/1807.10221>.
- [17] Enze Xie, Wenhai Wang, Zhiding Yu, Anima Anandkumar, Jose M Alvarez, and Ping Luo. Segformer: Simple and efficient design for semantic segmentation with transformers. In *Neural Information Processing Systems (NeurIPS)*, 2021.
- [18] Zongwei Zhou, Md Mahfuzur Rahman Siddiquee, Nima Tajbakhsh, and Jianming Liang. Unet++: Redesigning skip connections to exploit multiscale features in image segmentation. *IEEE Transactions on Medical Imaging*, 2019.